

Modeling with Deterministic Functions - Capturing Signal

Zack Treisman

Spring 2022



Philosophy

Our basic scenario: x is a **predictor** and y the **response**. We want

$$y = f(x) + \epsilon$$

where the model function $f(x)$ is what we call the **signal** and ϵ is the **noise**.

- ▶ Making good models means choosing an appropriate $f(x)$ and an appropriate ϵ .
- ▶ These choices are interdependent, but this week we'll focus more on signal than noise.

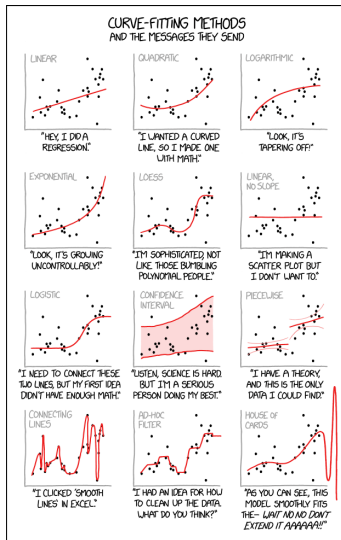


Figure 1: <https://xkcd.com/2048>

Model misuse is not a joke

United States Daily COVID-19 Deaths: Actual Data, IHME/UW Model Projections, & Cubic Fit.

Updated today (5/5/20), data through yesterday (5/4/20).

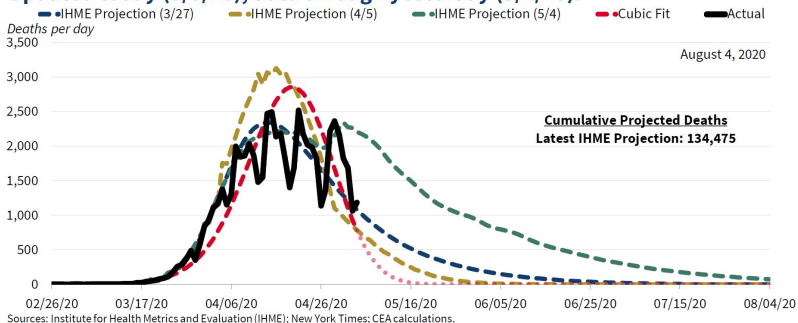


Figure 2:

<https://twitter.com/WhiteHouseCEA/status/1257680258364555264>

Reducible and irreducible error

If we had infinite knowledge, we could choose for our model function the expected value, or mean, of all y such that (x, y) is a possible data point. Write $\mu(x)$ for this expected value.

$$y = \mu(x) + \varepsilon$$

We call ε the **irreducible error** or **intrinsic variance**. It is the uncertainty that exists because of natural variation in the system described.


Any actual model function that we come up with will differ from this optimal function. Suppose we have a model function $f(x)$. We call the difference $f(x) - \mu(x)$ the **reducible error**. With a better f we can reduce the reducible error.

Bias and Variance

The reducible error can be broken down into two parts.

- ▶ **Bias** is that part of the reducible error that comes from a model function's inability to change when it needs to.
- ▶ **Variance** is that part of the reducible error that comes from a model function's flexibility to match particular data.

A model function with **high bias error** is said to **underfit** the data, and one with **high variance error** is an **overfit**. Whenever we are choosing a model, we must consider this **bias-variance tradeoff**.



Parametric vs Non-parametric models

- ▶ A **parametric** model function is one defined in terms of arithmetic and analytic functions, such as logarithms, polynomials, or anything else you might have encountered in a math class like Calculus. The coefficients, exponents *et cetera* defining the function are called the **parameters**.
 - ▶ Parametric models can be more interpretable and meaningful, especially if there are some deterministic processes in play.
 - ▶ In principle, any function can be used, but a great deal of the work in statistics is done by linear functions and a few others.
- ▶ A **non-parametric** model function is defined in some other way. The first example we will encounter and use is *local regression* or *loess*. Trees, random forests, neural networks and most other are also non-parametric models.
 - ▶ Non-parametric models often have better predictive power and can fit data more effectively.

Linear function models

If x and y are both numeric variables, then the simplest possible relationship between them is a linear relationship.

$$y = \beta_0 + \beta_1 x + \epsilon$$

- ▶ β_0 is the *intercept*, the value we predict for y when x is zero.
- ▶ β_1 is the *slope*, or predicted rate of change in y with respect to x . Often written $\frac{\Delta y}{\Delta x}$ or $\frac{dy}{dx}$.

The tremendous advantage of the linear function model over all others is its **simplicity**. A disadvantage is a tendency towards **high bias error**.

Note: This model is linear in the variable x **and** in the *parameters* β_i . The term “linear model” is frequently applied with both meanings, but the latter is more useful. The R command `lm` refers to linearity in the parameters.

Multiple predictors, categorical predictors

As we discussed last week, extending to multiple predictor variables or categorical predictors is relatively straightforward.

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \epsilon$$

- ▶ β_i is expected change in y as x_i changes and all else is fixed.
- ▶ If a predictor variable is categorical:
 - ▶ One level is set as the **reference** level of the variable.
 - ▶ For every other level of the variable, an **indicator variable** is defined, taking the value 1 when the variable has that level and 0 otherwise. My typical notation for an indicator variable is χ_i .
 - ▶ The coefficient β_i is the expected effect from observing level i instead of the reference level.
- ▶ If the x_i are correlated, these models can be unreliable.


Polynomial functions

Powers of x such as x^2 or x^3 **reduce bias** but **increase variance**.

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \epsilon$$

This is still a linear model even though it includes the x^2 term. The parameters β_i appear linearly.

```
ggplot(mtcars, aes(mpg, hp)) + geom_point(alpha=0.5) +  
  geom_smooth(method = lm, formula = y~x, se=F) +  
  geom_smooth(method = lm, formula = y~poly(x,2), se=F, color="red")
```



Power functions

A power function model looks like

$$y = ax^k + \epsilon$$

- ▶ k can be any number, not just a positive integer.
- ▶ This is **not a linear model** because of the parameter k .
- ▶ Power models can have explanatory meaning, for example if x and y have relevant dimensionality, like mass or area.

Power function models can be fit using linear model techniques by taking logarithms. Ignoring the error term for a moment:

$$\hat{y} = ax^k \quad \Longleftrightarrow \quad \log(\hat{y}) = \log(a) + k \log(x)$$

Back-transforms and error

Models fit on log-transformed variables can be exponentiated back to the original variables. How the error transforms can cause issues.

- ▶ Exponentiating the predicted mean of a log-transformed variable does **not** predict the untransformed mean.
 - ▶ When back-transforming, add half the variance in the residuals before exponentiating to recover the mean.
 - ▶ Diagnostics such as R^2 and p values apply to the transformed variables, not after back-transformation.
- ▶ Linear regression assumes that the error is additive. Exponentiation changes this addition into multiplication.

Suppose we fit a model:

$$\log(y) \sim N\left(\hat{\beta}_0 + \hat{\beta}_1 \log(x), \sigma\right).$$

Then the prediction for the mean of y is


$$e^{\hat{\beta}_0 + \hat{\beta}_1 \log(x) + \sigma^2/2} = e^{\hat{\beta}_0 + \sigma^2/2} x^{\hat{\beta}_1}$$

and the variance is dependent on x .

Kleiber's law

Mass and metabolic rate of mammals relate via a power law.

```
ggplot(ex0826, aes(Mass, Metab)) + geom_point() + # data in Sleuth3  
  scale_x_log10() + scale_y_log10() + geom_smooth(method = lm, se=F)
```



```
lm1 <- lm(log(Metab)~log(Mass), data = ex0826)  
lm1$coefficients; var(lm1$residuals)
```


```
## (Intercept)    log(Mass)  
##    5.6383307    0.7387436  
  
## [1] 0.2068395
```

$$\text{Metab} = e^{5.64+0.21/2} \times (\text{Mass})^{0.74} \times \epsilon$$

Kleiber's law

Mass and metabolic rate of mammals relate via a power law.

```
ggplot(ex0826, aes(Mass, Metab)) + geom_point() + # data in Sleuth3  
  scale_x_log10() + scale_y_log10() + geom_smooth(method = lm, se=F)
```



```
lm1 <- lm(log(Metab)~log(Mass), data = ex0826)  
lm1$coefficients; var(lm1$residuals)
```

```
## (Intercept)    log(Mass)  
##    5.6383307    0.7387436  
  
## [1] 0.2068395
```

$$\text{Metab} = e^{5.64+0.21/2} \times (\text{Mass})^{0.74} \times \epsilon$$

Exponential models

Exponential growth and decay are very common.

- ▶ $\frac{dy}{dx} = k \cdot y$
- ▶ Change in y is proportional (with coefficient k) to y itself.

The solution is

$$y = a e^{kx} + \epsilon, \quad a = \text{initial value}, \quad k = \text{growth or decay rate}$$

Growth if $k > 0$, decline if $k < 0$.

- ▶ Exponential growth is usually bad for extrapolation. Something else tends to take over.
- ▶ With a little algebra, exponential decay can be used to model convergence to any asymptote.

These models are not linear, but taking logs make them so:

$$\hat{y} = a e^{kx} \quad \Longleftrightarrow \quad \log(\hat{y}) = \log(a) + kx$$

Exponential decay example

Exponential decay can model decomposition, radioactive decay, temperature changes, and many other scenarios.


Simulate: $a = 10, k = 0.005 + 0.005\chi_{green} + 0.007\chi_{blue}, \sigma = 0.5$

$$\hat{mass} = 10 e^{(0.005 + 0.005\chi_{green} + 0.007\chi_{blue})age}, \quad \text{or}$$

$$\log(\hat{mass}) = \log 10 + (0.005 + 0.005\chi_{green} + 0.007\chi_{blue})age$$

Simulated data, see RMarkdown for code.

```
ggplot(decay, aes(age, mass, color=k))+geom_point()
```



Option 1: model using log-transformed mass

Take logs of all the mass values and then make a linear model.

```
lm2 <- lm(log(mass)~age*k, data=decay)
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	2.3243	0.0356	65.2844	0.0000
## age	-0.0054	0.0005	-11.7842	0.0000
## k0.01	-0.0019	0.0455	-0.0410	0.9674
## k0.012	-0.0078	0.0485	-0.1606	0.8728
## age:k0.01	-0.0046	0.0006	-7.5162	0.0000
## age:k0.012	-0.0069	0.0007	-9.7919	0.0000

Residual variance is 0.008. At $\alpha = 0.05$ this model says ...

- ▶ Significant intercept: the initial value for the mass for red points is not 1. (Its log is not 0.)
- ▶ Significant coefficient on age: $k \neq 0$ for the red points.
- ▶ No significant coefficients on indicator variables: initial mass does not appear to depend on color.
- ▶ Significant coefficients of the interaction terms for age and the indicator variables: decay rates differ by color.

Option 1 (continued)

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	2.3243	0.0356	65.2844	0.0000
## age	-0.0054	0.0005	-11.7842	0.0000
## k0.01	-0.0019	0.0455	-0.0410	0.9674
## k0.012	-0.0078	0.0485	-0.1606	0.8728
## age:k0.01	-0.0046	0.0006	-7.5162	0.0000
## age:k0.012	-0.0069	0.0007	-9.7919	0.0000

Plugging the significant coefficients in and simplifying, the model is:

$$\log(\text{mass}) = 2.3243 + (-0.0054 - 0.0046\chi_{\text{green}} - 0.0069\chi_{\text{blue}})\text{age} + \epsilon$$

Where $\epsilon \sim N(0, \sqrt{0.008})$. Transforming back to the original scale¹:

$$\text{mass} = 10.26048 \times e^{(-0.0054 - 0.0046\chi_{\text{green}} - 0.0069\chi_{\text{blue}})\text{age}} \times e^{\epsilon}$$

This is pretty close to the actual model, but there is a potential issue. The error has become multiplicative instead of additive, and it follows this “exponential normal” distribution.

¹using $10.26048 = e^{2.3243+0.008/2}$

Link functions

Many people who fit models to log transformed variables don't do the back-transformation. But this means that they aren't actually talking about the variables in the system they want to model.

Generalized linear models (GLMs) offer a solution:

Assume there is an invertible function $g(x)$ and define a model

$$y = g(\beta_0 + \beta_1 x) + \epsilon.$$

The function g^{-1} is called the **link**. For example, a GLM with a log link finds:

$$\log(\hat{y}) = \hat{\beta}_0 + \hat{\beta}_1 x$$

or equivalently

$$\hat{y} = e^{\hat{\beta}_0 + \hat{\beta}_1 x}.$$

Using GLMs and link functions means we can talk about our variables, not their transforms, and errors are homoscedastic.

Option 2: model using glm and log link

Specifying a link function for a GLM in R is done in the `family` argument (which is also used to describe the shape of the error). For normally distributed error, use `family=gaussian`.

```
lm2a <- glm(mass~age*k, family=gaussian(link=log), data=decay)
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	2.3232	0.0208	111.5830	0.0000
## age	-0.0054	0.0003	-16.6556	0.0000
## k0.01	-0.0082	0.0278	-0.2947	0.7689
## k0.012	-0.0362	0.0315	-1.1500	0.2531
## age:k0.01	-0.0045	0.0005	-8.9756	0.0000
## age:k0.012	-0.0062	0.0006	-9.5910	0.0000

The resulting model is

$$\text{mass} = 10.20829 \times e^{(-0.0054 - 0.0045\chi_{\text{green}} - 0.0062\chi_{\text{blue}})\text{age}} + \epsilon$$


where ϵ is normally distributed with standard deviation 0.4594.

Logistic models

The logistic function makes a transition from $y = 0$ to $y = 1$.

$$y = \frac{e^{a+bx}}{1 + e^{a+bx}}$$

```
curve(exp(-7+3*x)/(1+exp(-7+3*x)), from=0, to=4, ylim=c(0, 1), ylab="y")
```



- ▶ Mostly used for binary classification. (logistic regression)
- ▶ Also useful for populations with a carrying capacity.


Model with `link=logit` in `glm`. Default for `family=binomial`.

Other function models

Any function can potentially be used as a model function.

- ▶ Several pre-packaged link functions are available and work well with particular choices of error distribution.
- ▶ More complicated functions require more complicated code.

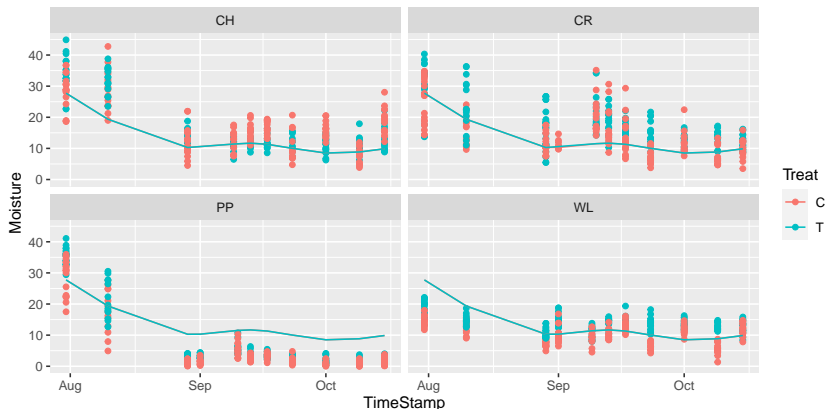
$$\mu_{titer} = a_g t e^{-b_g t}$$



Splines

Weather drives the response. A spline can help account for this.



```
SoilMoist <- read.csv("data/SoilMoisture_ALL.csv") # A. Cooper's data
SoilMoist$TimeStamp <- mdy(SoilMoist$TimeStamp) # from lubridate package
splineMod <- lm(Moisture~ns(TimeStamp,4), data=SoilMoist)
SoilMoist$pred <- predict(splineMod)
ggplot(SoilMoist, aes(TimeStamp,Moisture, color = Treat))+
  facet_wrap(~Site)+geom_point()+geom_line(aes(y=pred))
```



Nearest neighbor averaging (knn)

- Assume the actual expected value function doesn't change too quickly: $\mu(x - h) \approx \mu(x) \approx \mu(x + h)$.

Choose a positive integer k and define $f(x)$ as the average of y_i for the points (x_i, y_i) where $x - x_i$ is among the k smallest.




Left: knn with $k = 1$. Variance is high. Right: knn with $k = 9$.

Weighted averaging and local regression (loess)

Choose a distance and define $f(x)$ by linear regression using the data within that distance of x or weighted based on that distance.

Local Regression



Simulated data. Blue curve is the true signal, orange is the weighted local regression.

Loess in R

One way to do loess in R is to use `gam` and `lo`.

```
moistmod0 <- gam(Moisture~lo(TimeStamp), data=SoilMoist)
moistmod1 <- gam(Moisture~lo(TimeStamp)+Site, data=SoilMoist)
moistmod2 <- gam(Moisture~lo(TimeStamp)+Site+Treat, data=SoilMoist)
moistmod3 <- gam(Moisture~lo(TimeStamp)+Site*Treat, data=SoilMoist)
anova(moistmod0, moistmod1, moistmod2, moistmod3, test="F")
```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: Moisture ~ lo(TimeStamp)
```

```
## Model 2: Moisture ~ lo(TimeStamp) + Site
```

```
## Model 3: Moisture ~ lo(TimeStamp) + Site + Treat
```

```
## Model 4: Moisture ~ lo(TimeStamp) + Site * Treat
```

```
##      Resid. Df Resid. Dev Df Deviance      F      Pr(>F)
```

```
## 1      1362.8      56413
```

```
## 2      1359.8      38604  3  17809.1 225.919 < 2.2e-16 ***
```

```
## 3      1358.8      36577  1   2026.9  77.138 < 2.2e-16 ***
```

```
## 4      1355.8      35627  3    950.3  12.056 8.642e-08 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Using loess instead of a spline, we ask if Site, Treat, and their interaction are significant drivers of soil moisture after accounting for the common temporal variability. It appears that the answer is yes.

Acknowledgements

Some figures in this presentation are taken from “An Introduction to Statistical Learning, with applications in R” (Springer 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani

Other figures are created using code provided by Ben Bolker related to his text “Ecological Models and Data in R” (Princeton 2008)

References

- Bolker, Benjamin M. 2008. *Ecological Models and Data in r*. Princeton University Press.
- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2013. *An Introduction to Statistical Learning: With Applications in r*. Springer.
<https://faculty.marshall.usc.edu/gareth-james/ISL/>.